



•

•

•

•

•

•

•

•

•



•

•

•

•

•

•

•

•

•

•

•

•

-
-
-
-
-

•

■

■



•

•

•

•

•

•

•

•



```

import pyopencl as cl

print('\n' + '=' * 60 + '\nOpenCL Platforms and Devices')
# Print each platform on this computer
for platform in cl.get_platforms():
    print('=' * 60)
    print('Platform - Name: ' + platform.name)
    print('Platform - Vendor: ' + platform.vendor)
    print('Platform - Version: ' + platform.version)
    print('Platform - Profile: ' + platform.profile)
    # Print each device per-platform
    for device in platform.get_devices():
        print(' ' + '-' * 56)
        print(' Device - Name: ' + device.name)
        print(' Device - Type: ' +
cl.device_type.to_string(device.type))
        print(' Device - Max Clock Speed: {0}
Mhz'.format(device.max_clock_frequency))
        print(' Device - Compute Units:
{0}'.format(device.max_compute_units))
        print(' Device - Local Memory: {0:.0f}
KB'.format(device.local_mem_size/1024.0))
        print(' Device - Constant Memory: {0:.0f}
KB'.format(device.max_constant_buffer_size/1024.0))
        print(' Device - Global Memory: {0:.0f}
GB'.format(device.global_mem_size/1073741824.0))
        print(' Device - Max Buffer/Image Size: {0:.0f}

```

```
=====
OpenCL Platforms and Devices
=====
Platform - Name: Intel(R) OpenCL HD Graphics
Platform - Vendor: Intel(R) Corporation
Platform - Version: OpenCL 2.1
Platform - Profile: FULL_PROFILE
-----
Device - Name: Intel(R) Gen9 HD Graphics NEO
Device - Type: GPU
Device - Max Clock Speed: 1050 Mhz
Device - Compute Units: 24
Device - Local Memory: 64 KB
Device - Constant Memory: 4194296 KB
Device - Global Memory: 25 GB
Device - Max Buffer/Image Size: 4096 MB
Device - Max Work Group Size: 256
=====
```

•

•

•

•

•

■

■



•

•


```

import numpy as np
import pyopengl as cl

a_np = np.random.rand(50000).astype(np.float32)
b_np = np.random.rand(50000).astype(np.float32)

ctx = cl.create_some_context()
queue = cl.CommandQueue(ctx)

mf = cl.mem_flags
a_g = cl.Buffer(ctx, mf.READ_ONLY | mf.COPY_HOST_PTR,
hostbuf=a_np)
b_g = cl.Buffer(ctx, mf.READ_ONLY | mf.COPY_HOST_PTR,
hostbuf=b_np)

prg = cl.Program(ctx, """
__kernel void sum(
    __global const float *a_g, __global const float *b_g,
    __global float *res_g)
{
    int gid = get_global_id(0);
    res_g[gid] = a_g[gid] + b_g[gid];
}
""").build()

res_g = cl.Buffer(ctx, mf.WRITE_ONLY, a_np.nbytes)
prg.sum(queue, a_np.shape, None, a_g, b_g, res_g)

res_np = np.empty_like(a_np)
cl.enqueue_copy(queue, res_np, res_g)

# Check on CPU with Numpy:
print(res_np - (a_np + b_np))
print(np.linalg.norm(res_np - (a_np + b_np)))
assert np.allclose(res_np, a_np + b_np)

```

```
buf = cl . Buffer ( context , flags , size =0, hostbuf=None)
```

-
-
- -
 -
-
-

```
buf = cl . Buffer ( context , flags , size =0, hostbuf=None)
```

-
- -
- -
 -
- -

•

■

■

○


```
prg = cl.Program(context, src)
```

-
-

```
prg.build(options="", devices=None)  
kernel = prg.kernelname(queue, (Gx, Gy, Gz), (Lx, Ly, Lz), *args)
```

-
- -
 -
 -

```
kernel = prg.kernelname(queue, (Gx,Gy,Gz), (Lx,Ly,Lz), *args)
```

-
-
-

```

kernel = """
__kernel void sum(
    __global float* a,
    __global float* b,
    __global float* c,
    const unsigned int count)
{
    // your code here
}

"""

vector_size = 1024
context = cl.create_some_context()
queue = cl.CommandQueue(context)

h_a = np.random.rand(vector_size).astype(np.float32)
h_b = np.random.rand(vector_size).astype(np.float32)
h_c = np.empty(vector_size).astype(np.float32)

d_a = cl.Buffer(context, cl.mem_flags.READ_ONLY |
cl.mem_flags.COPY_HOST_PTR, hostbuf=h_a)
d_b = cl.Buffer(context, cl.mem_flags.READ_ONLY |
cl.mem_flags.COPY_HOST_PTR, hostbuf=h_b)
d_c = cl.Buffer(context, cl.mem_flags.WRITE_ONLY, h_c.nbytes)

program = cl.Program(context, kernel).build()
# Utiliser obligatoirement une variable pour utiliser une seule
instance (cf. doc)
prg = program.sum
prg.set_scalar_arg_dtypes([None, None, None, np.uint32])
prg(queue, h_a.shape, None, d_a, d_b, d_c, vector_size)
// ...

```

- **None**

-

```
printf("Hello from kernel #%d, got value: %d\n", global_id,  
values[global_id]);
```




•

■

■

■

■

•

•

•

•

```

import pyopencl as cl
import numpy

A = numpy.random.rand(1000).astype(numpy.float32)
B = numpy.random.rand(1000).astype(numpy.float32)
C = numpy.empty_like(A)

ctx = cl.Context()
queue = cl.CommandQueue(ctx)

A_buf = cl.Buffer(ctx, cl.mem_flags.READ_ONLY |
cl.mem_flags.COPY_HOST_PTR, hostbuf=A)
B_buf = cl.Buffer(ctx, cl.mem_flags.READ_ONLY |
cl.mem_flags.COPY_HOST_PTR, hostbuf=B)
C_buf = cl.Buffer(ctx, cl.mem_flags.WRITE_ONLY, A.nbytes)
prg = cl.Program(ctx, """
__kernel
void sum(__global const float* a, __global const float* b,
__global float* c){
    int i = get_global_id(0);
    c[i] = a[i] + b[i];
}
""").build()
prg = cl.Program(ctx, kernel).build()
prg.sum(queue, A.shape, A_buf, B_buf, C_buf)
cl.enqueue_copy(queue, C, C_buf)

```


•

■

■

•

■

■

•

■

```
void
mul(const int n,
    const float *a,
    const float *b,
    float *c)
{
    int i;
    for (i = 0; i < n; i++)
        c[i] = a[i] * b[i];
}
```

```
__kernel void
mul(__global const float *a,
    __global const float *b,
    __global float *c)
{
    int id = get_global_id(0);
    c[id] = a[id] * b[id];
}
```

•

•

•

•

-
-
-
-



•

•

•

■

•

■

•

■

•

•

■

■

○

○

○

•

•

-
-
-
-

-
-
-
-



•

■

■

■

•

■

■

■

•

■

■

■

•

■

■



•

•

•

•

•



-
-
-
-

•

•

•

•

•

•

•

•



•

■

•



2

2



-

-

```
barrier(CLK_LOCAL_MEM_FENCE)
```

```
barrier(CLK_GLOBAL_MEM_FENCE)
```

-

-

•

■

■

■

■

•

■

■

■

■



```
void mat_mul(int N, float *A, float *B, float *C) {  
    int i, j, k;  
    for (i = 0; i < N; i++) {  
        for (j = 0; j < N; j++) {  
            C[i*N+j] = 0.0f;  
            for (k = 0; k < N; k++) {  
                C[i*N+j] += A[i*N+k] * B[k*N+j];  
            }  
        }  
    }  
}
```

```
1  __kernel void mat_mul(const int N, __global float *A,
   __global float *B, __global float *C) {
2      int i, j, k;
3      // Les boucles sont transformées en calculs parallèles
4      i = get_global_id(0);
5      j = get_global_id(1);
6      // il ne reste que le produit scalaire
7      for (k = 0; k < N; k++) {
8          C[i*N+j] += A[i*N+k] * B[k*N+j];
9      }
10 }
```

```
1  __kernel void mmul(const int N,__global float *A,__global
   float *B,
2  __global float *C) {
3      int k;
4      int i = get_global_id(0);
5      int j = get_global_id(1);
6      float tmp = 0.0f;
7      for (k = 0; k < N; k++)
8          tmp += A[i*N+k]*B[k*N+j];
9      C[i*N+j] += tmp;
10 }
```

•

•



```
__kernel void mmul(const int N, __global float *A, __global
float *B,
__global float *C) {
    int j, k;
    int i = get_global_id(0);
    float tmp;
    for (j = 0; j < N; j++) {
        tmp = 0.0f;
        for (k = 0; k < N; k++)
            tmp += A[i*N+k]*B[k*N+j];
        C[i*N+j] = tmp;
    }
}
```


•

•



```

1  __kernel void mmul(const int N, __global float *A, __global
   float *B,
2      __global float *C) {
3      int k;
4      int j;
5      int i = get_global_id(0);
6      float tmp;
7      float Awrk[1024]; // Mémoire privée
8
9      // copie de la ligne de A en memoire privée
10     for (k = 0; k < N; k++) {
11         Awrk[k] = A[i * N + k];
12     }
13
14     // Multiplication d'une ligne entière
15     for (j = 0; j < N; j++) {
16         tmp = 0.0f;
17         for (k = 0; k < N; k++) {
18             tmp += Awrk[k] * B[k * N + j];
19         }
20         C[i * N + j] = tmp;
21     }
22 }

```

•

•

•

•


```

1  __kernel void mmul(const int N, __global float *A, __global
   float *B,
2      __global float *C, __local float *Bwrk)
   {
3      int k, j;
4      int i = get_global_id(0);
5      int iloc = get_local_id(0);
6      int nloc = get_local_size(0);
7
8      // Copy the row of A into private memory.
9      float Awrk[256];
10     float tmp;
11
12     for (k = 0; k < N; k++)
13         Awrk[k] = A[i * N + k];
14
15     for (j = 0; j < N; j++) {
16         // Initialisation de la colonne
17         for (k = iloc; k < N; k += nloc)
18             Bwrk[k] = B[k * N + j];
19
20     barrier(CLK_LOCAL_MEM_FENCE);
21
22     tmp = 0.0f;
23     for (k = 0; k < N; k++) {
24         tmp += Awrk[k] * Bwrk[k];
25     }
26     C[i * N + j] = tmp;
27
28     barrier(CLK_LOCAL_MEM_FENCE);
29 }
30 }

```

•

•

•

•

•

•

•

•

```
atomic_add, atomic_sub, atomic_xchg, atomic_cmpxchg,  
atomic_inc, atomic_and,  
atomic_or, atomic_xor, atomic_dec, atomic_max, atomic_min
```

```
__local int n;  
if (get_local_id (0) == 0) {  
    n = 0;  
}  
barrier (CLK_LOCAL_MEM_FENCE);  
//...  
atomic_inc (&n);  
//...  
}
```


•

•

•

```
import pyopencl as cl
import numpy as np

context = cl.create_some_context()
queue = cl.CommandQueue(context)

# 2 pyopencl arrays initialisées aléatoirement
a = cl.array.to_device(queue,
np.random.rand(50000).astype(np.float32))
b = cl.array.to_device(queue,
np.random.rand(50000).astype(np.float32))
# et le résultat
c = cl.array.empty_like(a)

program = cl.Program(context, """
__kernel void sum(__global const float *a, __global const float
*b, __global float *c)
{
    int i = get_global_id(0);
    c[i] = a[i] + b[i];
}""").build()

# Place le programme dans la file (queue) et récupère le
résultat dans c
program.sum(queue, a.shape, None, a.data, b.data, c.data)
```

•

•

•

$y_i = f_i(x_i)$ \text{, où :}

- f_i
- $f_i = \dots = f_N = f$
- $i \in \{1, \dots, N\}$



```
1  from pyopencl.elementwise import ElementwiseKernel
2
3  n = 10
4  a_np = np.random.randn(n).astype(np.float32)
5  b_np = np.random.randn(n).astype(np.float32)
6
7  ctx = cl.create_some_context()
8  queue = cl.CommandQueue(ctx)
9
10 a_g = cl.array.to_device(queue, a_np)
11 b_g = cl.array.to_device(queue, b_np)
12
13 lin_comb = ElementwiseKernel(ctx,
14     "float k1, float *a_g, float k2, float *b_g, float
15     *res_g",
16     "res_g[i] = k1 * a_g[i] + k2 * b_g[i]",
17     "lin_comb"
18 )
19 res_g = cl.array.empty_like(a_g)
20 lin_comb(2, a_g, 3, b_g, res_g)
```

$$y=f(\cdots f(f(x_1,x_2),x_3),\ldots,x_N)$$

?

?

ReductionKernel

-
-
-
-
-


```
1 a = pyopencl.array.arange(queue, 400, dtype=numpy.float32)
2 b = pyopencl.array.arange(queue, 400, dtype=numpy.float32)
3
4 krnl = ReductionKernel(ctx, numpy.float32, neutral="0",
5     reduce_expr="a+b", map_expr="x[i]*y[i]",
6     arguments="__global float *x, __global float *y")
7
8 prod = krnl(a, b).get() # recupère le numpy array
```

-
-

•

■

○

```
int4 v_iA = (int4)(7, -3, -2, 5);  
// ou int4 v_iA = {7, -3, -2, 5};  
  
int4 v_iB = (int4)(1, 2, 3, 4);  
int4 v_iC = v_iA + v_iB;
```



```
float4 pos = (float4)(4.0f, 3.0f, 2.0f, 1.0f);  
float4 reverse = pos.wzyx; //reverse = (1.0f, 2.0f, 3.0f, 4.0f)  
float4 duplicate = pos.xxyy; //duplicate = (4.0f, 4.0f, 3.0f,  
3.0f)
```

[]

```
int4 a;  
a.s0 = 2; // ou a[0] = 2
```

```
int2 a=(int2)(1,2);  
int2 b=(int2)(3,4);  
  
bool sup = all(a<b); // true
```

```
int4 i;  
*uint4 u = (uint4)i; // compile error  
  
// OK si le type de base est le même  
float4 f = 1.0f;  
float4 va = (float4)f; // va is a float4 vector  
// with elements ( f, f, f, f )  
  
// ou avec un scalaire  
float deux = 2.0f;  
int2 vc = (int2)deux; // vc is an int2 vector with elements  
// ( (int)deux, (int)deux )
```

```
uchar4 u;  
int4 c = convert_int4(u);  
  
float4 f = (float4)(-5.0f, 254.5f, 254.6f, 1.2e9f);  
uchar4 c = convert_uchar4_sat_rte(f); // c contient :  
// ((uchar)0, (uchar)254,(uchar)255,(uchar)255)
```

- sat
 - rte
-

pyopencl

```
ImageFormat([channel_order, channel_type])
```

```
Image(context, flags, format, shape=None, pitches=None,  
hostbuf=None, is_array=False, buffer=None)
```

-

-

- ~~image2d_t img[2]~~

- image2d_array_t

- struct

- image1d_t image2d_t image3d_t
- __read_only
- __write_only
- __read_write
-

- - read_imageui, read_imagei, read_imagef
write_imageui, write_imagei, write_imagef
 - float4, int4, uint4
 - cl_image_format
 -
 -

CL_SNORM_INT8, CL_UNORM_INT8, CL_SNORM_INT16,
CL_UNORM_INT16, CL_UNSIGNED_INT8, CL_UNSIGNED_INT16,
CL_UNSIGNED_INT32, CL_SIGNED_INT8, CL_SIGNED_INT16,
CL_SIGNED_INT32, CL_HALF_FLOAT, CL_FLOAT

```
numpy.asarray(...,order='F')  
ary.T.copy()
```

```
import pyopengl as cl
from PIL import Image

im = Image.open("image.png")
buffer = im.tobytes()

clImageFormat =
cl.ImageFormat(cl.channel_order.RGBA,cl.channel_type.UNORM_INT8
)

clImage = cl.Image(context,
                    cl.mem_flags.READ_ONLY |
cl.mem_flags.COPY_HOST_PTR,
                    clImageFormat,
                    im.size,
                    None,
                    buffer
                    )
```

```
from imageio import imread, imsave

#Read in image
im = imread('image.png').astype(np.float32) # type au choix...
```



```
__kernel void copie(__read_only image2d_t src,
__write_only image2d_t dest,
const sampler_t sampler,
const float sat)
{
    int row_id = get_global_id(0);
    int col_id = get_global_id(1);

    int2 coords;
    coords.y = row_id;
    coords.x = col_id;

    float4 pixel = read_imagef(src,sampler,coords);

    // simple copie des trois premiers plans
    write_imagef(dest,coords,pixel);
}
```

read_image

sampler_t

-

CLK_NORMALIZED_COORDS_TRUE

CLK_NORMALIZED_COORDS_FALSE



• read_image



- CLK_ADDRESS_MIRRORED_REPEAT CLK_ADDRESS_REPEAT
- CLK_ADDRESS_NONE

- filter Mode
 - CLK_FILTER_NEAREST CLK_FILTER_LINEAR
- - _____ __read_write

```
float4 read_imagef(image2d_t image,sampler_t sampler, float2
coord)

float4 read_imagef(image2d_t image,sampler_t sampler,int2
coord)

uint4 read_imageui(image2d_t image, sampler_t sampler,int2
coord)

int4 read_imagei(image2d_t image, sampler_t sampler,int2 coord)

...
```

float,int,uint
float2

```
void write_imagef(image2d_t image, int2 coord, float4 color)
void write_imagei(image2d_t image, int2 coord, int4 color)
...
```

•

•
